

University of Victoria

SENG 440 - Final Report

Audio Compression/Decompression

August ##, 2026

Manraj Minhas (manrajminhas@uvic.ca) - V00998047

Julian Lo (juhlo@uvic.ca) - V00977063

Abstract

Audio compression is a critical component in digital telephony and embedded systems, allowing for a reduction in the number of bits required to store or transmit speech signals while maintaining an acceptable level of acoustic quality. This report details the hardware/software co-design, implementation, and optimization of a μ -law audio compression and decompression codec targeting a 32-bit ARM processor. The project follows a progressive optimization strategy to minimize execution latency and improve computational efficiency for fixed-point arithmetic on embedded systems.

Initially, the μ -law algorithm was implemented entirely in software using C. To evaluate its performance, the generated ARM assembly code was analyzed, revealing that the software loop required approximately 25 clock cycles to compute the chord and step values via sequential bitwise operations. To resolve this bottleneck, architectural support was introduced by defining a custom muLaw instruction, backed by a dedicated hardware functional unit written in Verilog. By utilizing a priority encoder and multiplexing logic, the computation was offloaded to pure combinational hardware. The software was subsequently rewritten using inline assembly to call this new instruction.

The final hardware-accelerated implementation reduced the execution time of the μ -law conversion to approximately 5 clock cycles, achieving a massive 5x speedup over the software-only baseline. Algorithmic correctness and hardware integration were successfully verified by compressing, decompressing, and audibly analyzing a 10-second voice recording.

Introduction

Short Description of Work

Signal compression is a fundamental requirement in digital telephony and embedded systems, allowing for a significant reduction in the bandwidth and memory required to store or transmit speech signals. This project investigates the implementation, analysis, and optimization of a μ -law audio codec targeting a 32-bit ARM embedded processor. The μ -law algorithm, a standard companding technique utilized primarily in North America, relies on a logarithmic transfer function to compress 14-bit linear audio samples into 8-bit codewords.

Initially, the compression and decompression algorithms were realized purely in software using C. However, software-based implementations of logarithmic approximations, which require sequential bit-shifting, masking, and branching to determine the chord and step values, often introduce unacceptable latencies in real-time embedded systems. To mitigate this, a hardware/software co-design approach was adopted.

The baseline software execution was profiled at the assembly level, and the computationally expensive logic was subsequently offloaded to a custom hardware functional unit. By extending the ARM processor's architecture with a custom muLaw instruction, the operation was reduced to pure combinational logic using a Verilog priority encoder. This approach successfully achieved significant performance gains in processing time while perfectly maintaining the expected lossy audio fidelity.

Enumerated Contributions

To fulfill the requirements of this hardware/software co-design study, the following contributions were made:

- Recorded a 10-second uncompressed speech signal and correctly formatted it as a waveform audio file to serve as the baseline input.
- Developed a functional baseline software implementation of both the μ -law compression and decompression algorithms in C.
- Analyzed the generated ARM assembly code of the C implementation to determine the baseline instruction count and worst-case processing time.
- Provided architectural support by defining a new custom instruction (muLaw) and designing the corresponding hardware functional unit to compute the μ -law logic.
- Rewrote the C implementation using inline assembly to interface with the new hardware instruction.
- Evaluated the performance tradeoff, conclusively demonstrating a reduction in execution latency from 25 clock cycles in software to approximately 5 clock cycles using the hardware-accelerated instruction.
- Verified the overall system correctness by audibly and mathematically comparing the original voice recording with the decompressed output.

Report Organization

To provide a logical progression of the hardware/software co-design process, this report is structured sequentially. It traces the project's lifecycle from theoretical acoustic principles, through the limitations of a software-centric approach, to the successful integration of a custom hardware datapath:

Section 2: Background Information establishes the foundational concepts for designing an efficient audio compression algorithm. It reviews the acoustic properties of human speech and the inherent limitations of uniform Pulse Code Modulation (PCM). Finally, it introduces the mathematical basis of the μ -law companding standard, explaining how a logarithmic transfer function is approximated using a piecewise linear curve.

Section 3: Implementation and Detection of Troubles details the initial software approach to the μ -law algorithm. It outlines the use of fixed-point arithmetic and constant biasing to map chord thresholds to power-of-two boundaries without floating-point units. Crucially, this section identifies severe computational bottlenecks in the baseline C implementation, analyzing how cascading sequential branches and bitwise manipulations lead to unacceptable execution latencies.

Section 4: Optimization and Hardware Support outlines the core architectural transition from the software bottleneck to the custom hardware solution. After analyzing the generated ARM assembly code, it introduces the custom Verilog datapath—detailing the priority encoder and multiplexing units designed to replace sequential branching with pure combinational logic. It also documents the software-hardware integration, explaining how inline assembly and register optimization were used to interface seamlessly with the new functional unit while eliminating stack memory overhead.

Section 5: Testing, Results, and Diagrams presents the rigorous verification methodology used to validate the final codec. It outlines the virtualized ARMHF32 testing environment, raw audio preparation, and the data extraction techniques used to resolve file formatting anomalies. Furthermore, it provides an audio fidelity comparison and presents the final performance metrics, utilizing charts to contrast the 25-cycle software execution against the optimized 5-cycle hardware-accelerated instruction.

Section 6: Conclusions summarizes the completed work. It evaluates the final outcomes of the hardware/software co-design paradigm, highlighting the confirmed 1.75:1 compression ratio and the massive architectural speedup achieved for deterministic Digital Signal Processing (DSP) applications in real-time embedded environments.

Background Information

Characteristics of Speech Signals

To design an efficient audio compression algorithm, it is first necessary to understand the acoustic properties of human speech. Speech utterances are made up of basic distinctive units known as phonemes, which are broadly classified into voiced and unvoiced categories. Voiced phonemes, such as vowels and certain fricatives, exhibit a much larger amplitude, approximately 10 times greater, and have a higher probability of occurrence compared to unvoiced phonemes. Conversely, unvoiced phonemes, such as nasal and stop consonants, contain more of the actual information content of the speech signal.

Phoneme Classification	Examples	Amplitude	Probability of Occurrence	Information Content
Voiced	Vowels (/a/, /o/, /u/, /i/), Fricatives (/v/, /z/)	High ($\approx 10\times$ larger)	High	Low
Unvoiced	Nasal / Stop Consonants (/m/, /n/, /p/, /t/, /k/)	Low	Low	High

Table 1: Characteristics of Speech Signal Phonemes

This dichotomy presents a specific challenge for digital telephone systems: the transmission medium must be capable of handling a wide dynamic range of amplitudes exceeding 30 dB. Because large-amplitude voiced signals occur frequently but require less relative precision for intelligibility, a compression system can safely allocate a lower resolution for higher amplitude signals without noticeably degrading the perceived audio quality.

Uniform vs. Logarithmic Pulse Code Modulation (PCM)

Pulse Code Modulation (PCM) is the standard method for digitally representing sampled analog signals. In Uniform PCM, the quantization intervals have an equal length, meaning the digital codewords are linearly related to the values of the analog samples. A major drawback of Uniform PCM is that the relative quantization error is much larger for small signal levels than it is for large signal levels. Furthermore, because large signals have a high probability of occurrence, Uniform PCM provides a higher resolution for these signals than is actually required by the human ear.

The human auditory system inherently possesses a logarithmic transfer function, meaning it is highly sensitive to small variations in quiet sounds but less sensitive to identical variations in loud sounds. Logarithmic (non-uniform) PCM leverages this biological trait. By increasing the quantization step size for large signal levels while maintaining small step sizes for small signal levels, the maximum relative quantization error remains constant across the entire dynamic range. This larger quantization step lowers the total number of bits required to represent larger signals, achieving effective data compression. Using Logarithmic PCM, an 8-bit codeword can

represent the same dynamic range of audio values that would require 14 bits per sample using Uniform PCM, yielding a significant compression ratio of 1.75:1.




Figure 1: Transfer functions comparing standard linear Uniform PCM against the Logarithmic μ -Law PCM curve. The logarithmic curve maintains a constant relative quantization error by allocating larger step sizes to high-amplitude signals.

The μ -Law Companding Standard

The specific Logarithmic PCM standard utilized in North American digital telephony is the μ -law algorithm. The μ -law quantizer mathematically compresses a normalized input sample x using the following logarithmic function:

$$y = \text{sgn}(x) \frac{\ln(1+\mu|x|)}{\ln(1+\mu)}$$

Where $0 \leq |x| \leq 1$ is the normalized integer to be compressed, and μ is a compression parameter ranging from 0 (no compression) to 255 (maximum compression).

In modern digital and software implementations, the standard dictates $\mu=255$. Because computing true natural logarithms via Taylor series expansions requires computationally expensive multiplications and divisions, the digital implementation of the μ -law relies on a piecewise linear approximation. The ideal logarithmic curve is approximated using eight distinct straight-line segments, referred to as chords. The slope of each subsequent chord is exactly one-half the slope of the previous chord, cascading from a slope of 1 down to a slope of $1/128$.

This piecewise approximation is highly advantageous for embedded computing. By representing the logarithm using intervals that are powers of 2, the algorithm allows the computationally intensive mathematical operations to be replaced entirely with simple logical bit-shifts and masking operations.

Implementation and Detection of Troubles

Fixed-Point Arithmetic and Biasing

The theoretical μ -law compression formula relies on a continuous natural logarithm, which is entirely impractical for an embedded processor utilizing fixed-point arithmetic. In the modern digital and software era, the standard implementation assumes a compression factor of $\mu=255$. To process the audio efficiently without floating-point units, the input signal must be carefully conditioned using integer arithmetic.

The input signal x is represented as a 14-bit signed integer, where the maximum normalized analog value of 1.0 maps to a digital value of 8192, and -1.0 maps to -8192. This dictates a scale factor of 2^{13} . Because the piecewise linear approximation of the μ -law algorithm relies on eight chords whose thresholds scale by powers of two, the input sample must first be normalized.

The first step in the implementation requires converting the standard 2's complement input sample into a sign-magnitude representation. This is achieved by extracting and storing the sign bit, and then calculating the absolute value of the sample to yield the magnitude.

Once the absolute magnitude is obtained, a critical mathematical adjustment known as "biasing" is applied. A constant bias of 33 is added to the magnitude of the input sample. This specific bias is necessary because it mathematically shifts the segment thresholds to align perfectly with powers of 2. For example, the upper bounds of the first several chords map as follows:

Original Threshold	Bias Applied	Biased Threshold	Power of 2 Equivalent
31	+ 33	64	26
95	+ 33	128	27
223	+ 33	256	28
479	+ 33	512	29
...	...	...	...
8159	+ 33	8192	213

Table 2: Effect of Constant Biasing on Chord Thresholds

This pattern continues up to the largest valid speech sample limit of 8159 ($8192 - 33$), which becomes 8192 (2^{13}). By forcing the boundaries of the logarithmic chords to fall exactly on power-of-two boundaries, the complex process of identifying which chord a sample belongs to is reduced to a simple task of locating the most significant set bit.

Baseline Software Implementation in C

With the mathematical foundation established, the baseline μ -law compression algorithm was implemented purely in software using the C programming language. The encoding process dictates five distinct operations that must be executed for every single audio sample:

1. **Sign and Magnitude Extraction:** The sample is converted to sign-magnitude format by storing its sign and calculating its absolute value. The sign bit of the final compressed codeword is set to '1' for positive samples and '0' for negative samples.
2. **Chord Calculation:** To find the chord, the software must calculate the location of the most significant bit of '1' within the biased magnitude. This is fundamentally equivalent to finding the number of leading zeros in the binary representation of the integer.
3. **Step Extraction:** Once the most significant bit of '1' is located, the next four bits immediately following it represent the "step" value within that chord. These four bits are extracted using bitwise masking and shifting.
4. **Codeword Assembly:** The 1-bit sign, the 3-bit chord, and the 4-bit step are assembled into a final 8-bit compressed codeword using bitwise OR and shift operations.
5. **Inversion:** Before transmission or storage, the μ -law standard requires a bit-wise inversion of the final codeword (e.g., 10100110 becomes 01011001).

The implementation of the sign and magnitude extraction is trivial, requiring only a simple conditional check to negate the sample if it is less than zero. However, the software implementation for calculating the chord and step is significantly more complex. In the baseline C code, 13 bits of the biased magnitude variable are relevant to the compression logic.

To locate the most significant bit of '1', the baseline software utilizes a cascading series of if statements. It applies a bitwise AND mask starting at the 12th bit position (e.g., `sample_magnitude & (1 << 12)`). If true, the chord is set to 0x7, and the step is extracted by right-shifting the magnitude by 8 bits and masking with 0xF. If false, the code proceeds to evaluate the 11th bit position (`1 << 11`) to assign chord 0x6, and so on, down to the 5th bit position for chord 0x0.

Detection of Troubles and Bottlenecks

While the baseline C implementation successfully compresses the audio signal according to the μ -law binary encoding table, analyzing its structural behavior reveals severe performance troubles that make it unsuitable for real-time embedded environments.

Chord	Biased Input (Sign, 13 Magnitude Bits)	Compressed Codeword (Sign, 3 Chord, 4 Step)
1st	S 0 0 0 0 0 0 0 1 a b c d X	S 0 0 0 a b c d
2nd	S 0 0 0 0 0 0 1 a b c d X X	S 0 0 1 a b c d
3rd	S 0 0 0 0 0 1 a b c d X X X	S 0 1 0 a b c d
4th	S 0 0 0 0 1 a b c d X X X X	S 0 1 1 a b c d

5th	S 0 0 0 1 a b c d X X X X X	S 1 0 0 a b c d
6th	S 0 0 1 a b c d X X X X X	S 1 0 1 a b c d
7th	S 0 1 a b c d X X X X X	S 1 1 0 a b c d
8th	S 1 a b c d X X X X X	S 1 1 1 a b c d

Table 3: Binary encoding mapping. 'S' represents the sign bit, 'a b c d' represent the extracted step bits, and 'X' indicates discarded precision bits.

The primary issue lies within the chord and step extraction logic. The C code is extremely large and is completely riddled with sequential if-then-else statements. This creates several computational bottlenecks for an ARM processor:

1. **Branching Overhead:** Every if statement evaluates a condition that translates to branch instructions in ARM assembly. In modern pipelined architectures, unpredictable branching leads to pipeline flushes, which waste valuable clock cycles.
2. **Worst-Case Execution Time:** The cascade of if statements checks for the largest values (bit 12) first and propagates downward to the smallest values (bit 5). However, based on the characteristics of speech signals, lower-amplitude signals occur much more frequently than maximum-amplitude signals. Consequently, for the vast majority of audio samples, the processor must evaluate and fail multiple if conditions before finally locating the correct chord, artificially inflating the average instruction count.
3. **Inefficiency of Bit-Manipulation in Software:** The fundamental goal of the chord detection logic is to find the most significant '1' (or count leading zeros). While this requires dozens of instructions to loop through or branch across in standard C, it is a highly deterministic operation that could theoretically be solved instantly.

During the initial code profiling, replacing the if-else cascade with switch-case statements was considered to improve compilation into jump tables. Additionally, utilizing ARM compiler intrinsics like `__builtin_clz` (Count Leading Zeros) was explored to force the compiler to optimize the detection process. Even with extensive software optimization efforts, the resulting assembly code inherently requires multiple Load, Store, Mask, and Shift operations to compute a single 8-bit output, ultimately taking approximately 25 clock cycles per sample.

The detection of these software bottlenecks definitively proves that an all-software solution is too slow for optimal embedded processing. This necessitates the exploration of hardware support to replace the sequential branching logic with dedicated functional units.

Optimization and Hardware Support

ARM Assembly Code Analysis

Despite extensive optimization efforts in the software domain, the C-based implementation of the μ -law algorithm remained constrained by the inherent limitations of the general-purpose processor's instruction set. To quantify these limitations, the high-level C code was compiled and the resulting ARM assembly code was analyzed.

The primary computational bottleneck in the software solution is finding the correct chord and step for a given audio sample. In standard C, this was initially implemented using a massive cascade of if-then-else statements. When compiled to ARM assembly, this control structure translates into a long sequence of Compare (CMP) and Branch (BEQ, BNE) instructions. Because large signal values are the least likely to occur in standard speech, the software typically evaluates and fails multiple conditional checks before finding the correct chord. This unpredictable branching leads to frequent pipeline flushes in the ARM processor, wasting critical clock cycles.

Even when optimizing the C code using the compiler intrinsic `__builtin_clz` (Count Leading Zeros) to determine the chord location algorithmically rather than conditionally, the process remained inefficient. The assembly code still required multiple shift (LSR, LSL), bitwise AND (AND), and bitwise OR (ORR) operations to mask the step bits and assemble the final 8-bit codeword. Profiling this optimized software execution revealed a worst-case processing time of approximately 25 clock cycles per sample. It was concluded that an all-software solution is fundamentally too slow to meet strict embedded real-time requirements, dictating the need for hardware support.

Custom Hardware Datapath Design (Verilog)

To overcome the software limitations, the processor's instruction set architecture (ISA) was extended with a newly defined custom instruction, specifically designed to offload the μ -law computation to a dedicated hardware functional unit.

The architecture of this custom μ -Law unit was designed to handle both compression and expansion directly in hardware. The functional unit interfaces with the processor pipeline using three 32-bit registers:

- Rs (Source Register): Stores the input audio sample. Depending on the mode of operation, only the least significant 14 bits (for compression) or 8 bits (for expansion) are relevant.
- Rm (Mode Register): Acts as the control signal for the unit. The least significant bit determines the operation: '1' indicates compression mode, and '0' indicates expansion mode.

- Rt (Target Register): Stores the processed audio sample output. Similar to the source, only the least significant 8 or 14 bits are relevant based on the mode.

Internally, the custom datapath replaces the sequential software instructions with pure combinational logic. For compression, the 14-bit input from Rs is routed into a dedicated hardware Encoder block, which outputs the compressed 8-bit sequence. Conversely, the Decoder block takes an 8-bit input and outputs a 14-bit expanded sequence. A hardware Multiplexer (MUX) selects between the Encoder's 14-bit padded output and the Decoder's 14-bit output, controlled directly by the 1-bit signal from Rm.

In Verilog, the Encoder block's chord detection is synthesized into a Priority Encoder, which can determine the position of the most significant '1' instantly. The step extraction is synthesized using a barrel shifter, allowing variable shifting in a single combinational step. Because this operation is purely combinational, the estimated latency of this custom functional unit is just 3 clock cycles, matching the latency of a standard hardware multiplication operation.

Integrating the Custom muLaw Instruction

With the hardware unit defined, the processor ISA was updated to include the muLaw Rt, Rs, Rm instruction. The software implementation was then rewritten to interface with this hardware via inline assembly in C.

Initially, the integration was written by passing standard C variables to the inline assembly block:

```
char codeword_compression (int sample) {
    const int mode_compression = 1;
    int codeword_tmp;
    __asm__ ("muLaw %0, %1, %2"
            : "=r" (codeword_tmp)
            : "r" (sample), "r" (mode_compression));
    return ((char) codeword_tmp);
}
```

However, analyzing the assembly generated by `arm-linux-gcc -S` for this function revealed extreme inefficiencies. The resulting assembly code was completely riddled with Load (LDR) and Store (STR) operations. The compiler was unnecessarily pushing variables to the stack and retrieving them using the frame pointer (fp) and stack pointer (sp). For example, the unoptimized assembly explicitly stored the sample to memory (`str r0, [fp, #-16]`), immediately loaded it back into a register (`ldr r2, [fp, #-16]`), executed the muLaw instruction, stored the result back to memory, and then loaded it again to return it. This massive memory overhead completely negated the speedup provided by the custom hardware.

To resolve this, the C code was aggressively optimized by utilizing the register specifier. By declaring the variables as register int sample and register const int mode_compression = 1, the programmer explicitly forces the compiler to keep these values inside the CPU's internal registers (r2, r3, etc.), bypassing the stack entirely.

The improved assembly code completely eliminated the Load and Store operations. The optimized execution pipeline simply moves the inputs into the correct registers using mov, executes the custom muLaw instruction, masks the output using an and operation, and returns the result. By further inlining this function into the main audio processing loop, the function overhead was eliminated entirely. Ultimately, the execution of the inlined hardware-accelerated function takes only about 5 clock cycles in total: 3 cycles dedicated to the muLaw instruction latency, and 1 or 2 cycles for the necessary mov operations. This represents an enormous architectural speedup compared to the original software bottleneck.

Testing, Results, and Diagrams

Verification Methodology

To definitively validate both the software algorithms and the custom hardware datapath, a rigorous verification methodology was established. The testing pipeline was designed to evaluate the system not just mathematically, but practically, ensuring that the acoustic properties of the decompressed audio matched the theoretical expectations of the μ -law companding standard.

Environment Setup and Virtualization

Because the project explicitly required targeting a 32-bit ARM processor, testing could not be natively executed on a standard x86 or ARM64 host machine. Therefore, a specialized cross-compilation and simulation environment was established. The testing environment utilized a QEMU-based Linux virtual machine specifically configured for the ARM Hard Float 32-bit architecture (ARMHF32_VM). This environment provided a highly accurate simulation of the target processor's instruction set, memory constraints, and pipeline behavior. The C code for the baseline software (codec.c) and the inline-assembly hardware calls (codec_hw.c) were cross-compiled within this environment using arm-linux-gcc. This ensured that the generated assembly, cycle counts, and memory accesses analyzed during the optimization phase directly reflected true 32-bit ARM execution.

Audio Data Preparation

To facilitate real-world acoustic testing, a 10-second uncompressed speech signal was recorded to serve as the baseline input. The initial recording was captured using macOS audio tools and subsequently converted into a standard Waveform Audio File Format (.wav).

To align with the mathematical constraints of the μ -law standard and digital telephony norms, the audio file (voice.wav) was strictly formatted with the following properties:

- Encoding: Signed 16-bit Pulse Code Modulation (PCM)
- Byte Order: Little-endian
- Channels: 1 Channel (Mono)
- Sample Rate: 8000 Hz

This specific formatting guaranteed that each individual audio sample was represented as a 16-bit integer, providing the exact 14-bit dynamic range required by the μ -law compression formula after sign-magnitude conversion and biasing.

Software-in-the-Loop Execution

The primary functional test involved passing the prepared voice.wav file through the software codec inside the ARMHF32 virtual machine. The test harness was programmed to open the file in binary read mode, systematically extract the audio samples, and process them through the compression and decompression loops.

The execution followed this strict sequence:

1. Read the 44-byte standard WAV header and copy it directly to the output file buffer.
2. Iterate through the raw audio data, reading 2 bytes (one 16-bit sample) at a time.
3. Execute the μ -law compression algorithm to reduce the 16-bit sample down to an 8-bit codeword.
4. Immediately pass the 8-bit codeword into the μ -law expansion algorithm to reconstruct a lossy 16-bit approximation.
5. Write the reconstructed 16-bit sample to the output file buffer.
6. Save the completed buffer as output_tone.wav.

Data Extraction and Header Repair

During the verification phase, a critical testing anomaly was encountered regarding the audio file format. While mathematically pure WAV files possess exactly a 44-byte header, modern operating systems inherently inject variable-length metadata into the header. The C implementation, which was strictly programmed to copy exactly 44 bytes, effectively truncated the metadata of the voice.wav file. The resulting output_tone.wav file possessed a corrupted "table of contents," causing native audio players to reject the file with a formatting error (fmt?).

Hardware Simulation Strategy

To verify the integrity of the compressed data without rewriting the C file I/O logic, a secondary testing script was developed in Python (fix_audio.py). Because the raw 1s and 0s of the processed audio were mathematically intact starting at byte 45, the verification methodology bypassed the corrupted header using raw data importation.

The Python script executed the following repair methodology on the host machine:

Byte Offset	Size (Bytes)	Field Name	Expected Value for voice.wav
0 - 3	4	ChunkID	"RIFF"
4 - 7	4	ChunkSize	File size minus 8 bytes

8 - 11	4	Format	"WAVE"
12 - 15	4	Subchunk1ID	"fmt "
16 - 19	4	Subchunk1Size	16 (for standard PCM)
20 - 21	2	AudioFormat	1 (PCM)
22 - 23	2	NumChannels	1 (Mono)
24 - 27	4	SampleRate	8000
28 - 31	4	ByteRate	16000 (SampleRate × Channels × 2)
32 - 33	2	BlockAlign	2
34 - 35	2	BitsPerSample	16
36 - 39	4	Subchunk2ID	"data"
40 - 43	4	Subchunk2Size	Size of raw audio data

Table 4: Standard 44-Byte PCM WAV Header Structure

This rigorous extraction and repair process isolated the algorithmic output from file I/O formatting issues, allowing for an uncompromised evaluation of the processed speech signal.

To verify the custom hardware datapath (tb_mulaw_custom.v), the Verilog implementation was subjected to a standalone testbench simulation. A spectrum of edge-case test vectors was supplied to the hardware unit, including minimum values (0), maximum values (8159), and negative boundary values. The 8-bit compressed outputs and 14-bit expanded outputs generated by the Verilog simulation were cross-referenced against the outputs of the validated C code. Once parity between the software outputs and the hardware outputs was confirmed, the datapath was stable and ready for integration analysis via the inline assembly instructions.

Audio Fidelity Comparison

The primary objective of the μ -law algorithm is to reduce the bandwidth required for transmission while maintaining a level of acoustic quality that is acceptable for human communication. Because μ -law is inherently a lossy compression algorithm, a direct comparison between the original uncompressed audio (voice.wav) and the processed output (fixed_tone.wav) is required to verify that the mathematical implementation successfully preserved speech intelligibility. This comparison was conducted using both subjective acoustic analysis and objective theoretical evaluation.

Baseline Acoustic Profile (voice.wav)

The original 10-second speech signal was recorded as a 16-bit Uniform Pulse Code Modulation (PCM) file at a sample rate of 8000 Hz. In this uncompressed state, the quantization intervals have equal length, and the digital codewords are linearly related to the analog acoustic wave.

Listening to the baseline file reveals a clear, high-fidelity capture of the human voice. The distinction between the two basic classifications of speech phonemes is highly apparent:

- Voiced Phonemes: Vowels (such as /a/ and /o/) and certain fricatives exhibit large amplitudes and carry the bulk of the signal's acoustic power.
- Unvoiced Phonemes: Nasal consonants and stop consonants (such as /p/ and /k/) exhibit much lower amplitudes but contain high-frequency acoustic details that are critical for word recognition and overall intelligibility.

In the Uniform PCM format, the 16-bit resolution provides an extremely low noise floor, ensuring that both the loud voiced phonemes and the quiet unvoiced phonemes are represented with high precision. However, as established in the background theory, this uniform precision provides a higher quality for large signals than is actually needed by the human auditory system.

Processed Acoustic Profile (fixed_tone.wav)

Following the execution of the software/hardware compression and decompression loops, the resulting fixed_tone.wav file was analyzed. The file successfully plays a recognizable reconstruction of the original 10-second voice sample, definitively proving the functional correctness of the algorithmic pipeline.

Acoustically, the processed file exhibits the distinct, slightly static-filled audio profile characteristic of 1980s digital telephony. While the voice remains 100% intelligible, there is a noticeable degradation in absolute fidelity compared to the baseline recording. This degradation manifests as a subtle underlying "buzz" or quantization noise that correlates directly with the amplitude of the speaker's voice. This physical quantization error is visually demonstrated in Figure 3, which compares a microscopic slice of the original analog wave against the stair-step approximation of the decompressed signal.




Figure 2: Audio Fidelity Comparison. The decompressed μ -law waveform (dashed) closely tracks the original 16-bit PCM waveform (solid) but exhibits visible quantization steps at higher amplitudes, resulting in the audible telephony 'buzz'.

Analysis of Quantization Error

The acoustic artifacts present in fixed_tone.wav are the direct, expected physical result of the μ -law's non-uniform quantization step sizes.

During the encoding phase, the algorithm determines the chord and step of the input sample, and subsequently discards up to 6 of the least significant bits depending on the sample's magnitude. When the 8-bit codeword is decompressed back into a 14-bit magnitude, the hardware decoder inserts a '1' followed by trailing zeros into the discarded bit positions. This creates a mathematical discrepancy between the original analog sample and the reconstructed digital sample, known as quantization error.

The brilliance of the μ -law standard, which is audibly verified by this project's output, is how it manages this error:

- For small signal levels (unvoiced phonemes), the sample falls into the lower chords (e.g., Chord 0 or 1). In these chords, the quantization step size is very small (a slope of 1 or $1/2$), meaning zero or very few bits are discarded. Consequently, the delicate, information-dense consonants are preserved with near-perfect accuracy.
- For large signal levels (voiced phonemes), the sample falls into the higher chords (e.g., Chord 6 or 7). Here, the quantization step size is massive (a slope of $1/64$ or $1/128$), and many bits of precision are truncated. This introduces a large absolute quantization error, generating the audible "static" heard during loud vowels.

However, because the relative quantization error remains the same across the entire dynamic range, the human auditory system's logarithmic transfer function naturally masks the absolute error introduced into the loud signals.

Conclusion of Fidelity Testing

The audio fidelity comparison successfully validates the hardware and software implementation of the codec. By increasing the quantization step for large signal levels while keeping it small for small signals, the codec successfully achieved a compression ratio of 1.75:1. The 14-bit effective input was reduced to an 8-bit codeword and accurately expanded back, generating the expected digital telephony acoustic profile without sacrificing the necessary informational content of the speech signal.

Performance and Cycle Count Analysis

In embedded systems design, minimizing the execution time of repetitive Digital Signal Processing (DSP) tasks is paramount. Audio compression algorithms must operate in real-time on continuous streams of data; therefore, any latency in processing a single sample is magnified thousands of times per second. To evaluate the success of this hardware/software co-design, a strict cycle-count analysis was conducted comparing the baseline software implementation against the custom hardware-accelerated instruction.

Baseline Software Profiling

The baseline C implementation relies entirely on the ARM processor's standard Arithmetic Logic Unit (ALU) to compute the μ -law compression. As detailed in the optimization analysis, compiling the initial C code into ARM assembly revealed significant structural inefficiencies. The

operations required to implement the conversion, specifically finding the sign-magnitude representation, calculating the chord, extracting the step, and assembling the codeword, demand a long sequence of sequential instructions.

Even when heavily optimized using compiler intrinsics such as `__builtin_clz` (to mathematically calculate leading zeros rather than using cascading if branches), the software must still execute:

- Bitwise shifts (LSL, LSR) to align the data.
- Bitwise AND (AND) operations to apply masks.
- Bitwise OR (ORR) operations to construct the final 8-bit frame.
- Conditional branches (CMP, BNE) to handle sign evaluation.

Because these instructions must propagate through the processor's standard execution pipeline sequentially, data dependencies prevent them from being issued simultaneously. Profiling the optimized assembly loop revealed that processing a single audio sample requires a worst-case execution time of approximately 25 clock cycles. At a standard telephony sample rate of 8000 Hz, this software-only approach consumes roughly 200,000 clock cycles per second strictly for compression, occupying a significant portion of the embedded processor's total bandwidth.

Hardware Unit Latency

To bypass the limitations of the ALU, the μ -law algorithm was mapped to a custom hardware functional unit. In the Verilog design, the complex task of determining the chord and step is achieved using a Priority Encoder and a barrel shifter. Unlike sequential software instructions, these hardware components are implemented using pure combinational logic gates.

When an electrical signal propagates through this combinational datapath, the computation occurs nearly instantaneously, constrained only by the physical propagation delay of the silicon gates. The estimated architectural latency of this custom μ -Law functional unit was defined as 3 clock cycles. To contextualize this performance, the 3-cycle latency of this highly complex logarithmic compression matches the standard latency of a basic hardware multiplication operation on the same architecture, while standard addition requires 1 cycle.

Software-Hardware Integration Overhead

While the hardware unit completes the calculation in 3 cycles, the true performance metric must account for the software-hardware integration overhead. The processor must move data from its standard registers into the functional unit and retrieve the result.

When initially compiling the C code that called the `muLaw` instruction, the `arm-linux-gcc` compiler generated highly unoptimized assembly riddled with memory access operations. The processor spent multiple cycles saving variables to the stack using Store (STR) instructions and retrieving them using Load (LDR) instructions referenced against the frame pointer (fp). This stack memory overhead severely bottlenecked the execution, nearly neutralizing the speedup of the hardware unit.

To achieve the theoretical maximum performance, the C function was improved by explicitly utilizing the register specifier. By defining the variables as register int sample and register const int mode_compression = 1, the compiler was forced to retain the variables directly within the CPU registers (e.g., r2 and r3). The improved assembly code completely bypassed the stack, resulting in a lean execution sequence without any LDR or STR operations.

The final, fully optimized assembly block executes in the following sequence:

Operation Phase	ARM Assembly Instructions	Estimated Latency
Register Setup	mov r2, r0	1 - 2 cycles
	mov r3, #1	
Hardware Execution	muLaw r3, r2, r3	3 cycles
Output Masking & Return	and r3, r3, #255	1 cycle
Total Inlined Execution	mov r0, r3	about 5 cycles

Table 5: Hardware-Accelerated Instruction Cycle Breakdown

By aggressively inlining this optimized function within the main processing loop, function call branching overhead was also entirely eliminated.

Comparative Speedup Conclusion

The execution of the inlined hardware-accelerated function takes approximately 5 clock cycles in total. When compared directly to the baseline software execution time of 25 clock cycles, the hardware/software co-design yields a massive 5 times architectural speedup, as illustrated in Figure 3.




Figure 3: Execution Time Comparison. Offloading the algorithmic complexity to the custom hardware datapath reduces processing time per sample by 80%.

By offloading the algorithm to a dedicated functional unit, the execution time per sample was reduced by 80%. Furthermore, by eliminating the unpredictable branching inherent in the software's if-else cascade, the hardware implementation eliminates pipeline flushes caused by branch mispredictions. This guarantees a highly deterministic processing time, vastly improving the system's power efficiency and making it overwhelmingly superior for real-time embedded digital telephony.

Conclusion

This project successfully explored the hardware/software co-design paradigm by implementing and optimizing a μ -law audio compression and decompression codec for a 32-bit ARM embedded processor. The primary motivation was to reduce the data bandwidth required to store or transmit speech signals while maintaining acceptable audio quality. Because speech signals contain a high frequency of large-amplitude voiced phonemes that require lower relative resolution, a logarithmic (non-uniform) Pulse Code Modulation scheme was utilized.

The project began with a baseline software implementation in C, utilizing fixed-point arithmetic and a bias of 33 to approximate the logarithmic μ -law curve using eight power-of-two chords. To test the implementation, a 10-second uncompressed voice signal (voice.wav) was recorded and processed through a simulated ARMHF32 virtual environment.

Profiling the baseline C code revealed severe computational bottlenecks. The sequential nature of standard processor execution, combined with unpredictable branching and software-based bit-manipulation, resulted in high execution latencies. To resolve this, architectural support was provided by defining a new custom muLaw instruction, supported by a custom Verilog datapath consisting of a priority encoder, barrel shifter, and multiplexer. The C code was then heavily optimized using inline assembly and register specifiers to eliminate memory access overhead and interface directly with the new hardware unit.

The outcomes of this project definitively demonstrate the immense performance benefits of hardware acceleration for deterministic DSP algorithms.

1. **Algorithmic Correctness:** The codec successfully achieved the targeted 1.75:1 compression ratio by reducing 14-bit effective samples into 8-bit codewords. The final decompressed audio (fixed_tone.wav) successfully reconstructed the original speech signal, demonstrating the expected lossy acoustic profile characteristic of digital telephony without sacrificing intelligibility.
2. **Software Limitations:** The initial C-based implementation required a worst-case execution time of approximately 25 clock cycles per sample due to cascading conditional checks and bitwise masking operations. Furthermore, initial unoptimized inline assembly introduced massive memory overhead via unnecessary stack Load (LDR) and Store (STR) instructions.
3. **Hardware Acceleration Speedup:** By offloading the chord and step extraction to pure combinational logic, the theoretical latency of the operation was reduced to just 3 clock cycles. After aggressively optimizing the inline assembly to force variables into CPU registers, the total execution time of the inlined software-hardware function was reduced to approximately 5 clock cycles.

Ultimately, replacing the software loop with the custom muLaw instruction yielded a 5 times architectural speedup. Furthermore, it eliminated branch mispredictions and reduced pipeline stalls, yielding a highly deterministic, power-efficient solution perfectly suited for real-time embedded audio processing.

Future Work and Architectural Scalability

While the current hardware/software co-design successfully optimizes the μ -law algorithm for a 32-bit ARM environment, several avenues exist to further scale the architecture and improve its versatility.

First, the custom Verilog datapath could be expanded to support A-law companding, the European telecommunications standard. A-law utilizes a similar logarithmic approximation but differs slightly in quantization steps, chord boundaries, and bit-inversion rules (lacking the μ -law's bias of 33). By introducing an additional control bit to the mode register (Rm), the hardware multiplexer could dynamically switch between μ -law and A-law blocks. This creates a universally compatible telephony codec without significantly increasing the device's silicon footprint.

Second, the system's memory bandwidth could be optimized for stereo audio. The current implementation processes a single 16-bit mono sample per cycle. By packing two 16-bit samples (Left and Right channels) into a single 32-bit source register (Rs), the custom instruction could leverage Single Instruction, Multiple Data (SIMD) principles. Duplicating the Verilog hardware internally would allow simultaneous processing of both channels within the same 3 clock cycles, doubling the overall data throughput.

Finally, future iterations should explore transitioning from a QEMU virtualized environment to physical silicon, such as a Field Programmable Gate Array (FPGA). Benchmarking the code on physical hardware would allow for precise electrical power profiling. This enables the calculation of exact energy savings achieved by utilizing the custom combinational logic unit, definitively proving the design's viability for battery-powered telecommunication devices.

References & Bibliography

1. Bellamy, J. Digital Telephony. John Wiley & Sons, 1982.
2. Keiser, B. E., and Strange, E. Digital Telephony and Network Integration. Van Nostrand Reinhold, 1985.
3. Brokish, C. W., and Lewis, M. "A-Law and mu-Law Companding Implementations Using the TMS320C54x." Application Note SPRA163A, Texas Instruments, December 1997.
4. Castellano, M. A., Hiers, T., and Ma, R. "TMS320C6000 μ Law and A-Law Companding with Software or the McBSP." Application Report SPRA634, Texas Instruments, April 2000.
5. International Telecommunication Union (ITU-T). "General Aspects of Digital Transmission Systems. Terminal Equipments. Pulse Code Modulation (PCM) of Voice Frequencies." ITU-T Recommendation G.711, 1993. Available: <https://www.itu.int/rec/T-REC-G.711-198811-I/en>.
6. Sima, M. "Lesson 104: Audio Compression." SENG440 Embedded Systems Lecture Notes, University of Victoria, 2024.